

- If a static variable is declared inside a compound statement, only one copy exists.
- The backslashes at the end of each line are necessary to tell the preprocessor that the macro definition continues on the next line.

volatile tells the compiler that the value stored in memory may change at any time

x86 Assembly

BASIC

- ## CONDITION

- ## CODING

- ## CALLING CONVENTION

- Parameters are pushed from right to left
- For pointers and integers no more than 32 bits, the return value is placed in EAX.
- Integers and other non-floating-point types of more than 32 but no more than 64 bits are split between EDX (high bits) and EAX (low bits)
- EBX, ESI, and EDI are callee-saved

<code>%ip</code>	instruction pointer (program counter)	31 16 15 8 7 0			
<code>%eflags</code>	flags (condition codes and other things)	<table border="1"><tr><td style="width: 16px; height: 1em;"></td><td style="width: 8px; height: 1em; text-align: center;">AH</td><td style="width: 8px; height: 1em; text-align: center;">AL</td></tr></table>		AH	AL
	AH	AL			

```
leal LABEL, %edx == movl $LABEL, %edx
```

type	generated by	example	asynchronous	unexpected
interrupt	external device	packet arrived at network card	yes	yes
exception	invalid opcode or operand	divide by zero	no	yes
system call/trap	deliberate, via INT instruction	print character to console	no	no

- A critical section is a block of code that executes a set of operations that should be executed without interruption. The critical section occurs atomically with respect to the interrupt.
- On computers with only a single processor, the approach is simple enough: use interrupt masking at the boundaries of critical sections
- Make critical sections as short as possible.

us	unexpected
	yes
	yes
	no

INTERRUPT CHAINING: multiple handler can be triggered by one interrupt

us	unexpected
	yes
	yes
	no

0x00–0x1F	0x00	division error	defined by Intel
	⋮		
	0x02	NMI (non-maskable interrupt)	
	0x03	breakpoint (used by KGDB)	
	0x04	overflow	
	⋮		
	0x0B	segment not present	
	0x0C	stack segment fault	
	0x0D	general protection fault	
	0x0E	page fault	
	⋮		
0x20–0x27	0x20	IRQ0 — timer chip	example of possible settings
	0x21	IRQ1 — keyboard	
	0x22	IRQ2 — (cascade to slave)	
	0x23	IRQ3	
	0x24	IRQ4 — serial port (KGDB)	
	0x25	IRQ5	
	0x26	IRQ6	
0x28–0x2F	0x27	IRQ7	
	0x28	IRQ8 — real time clock	slave 8259 PIC
	0x29	IRQ9	
	0x2A	IRQ10	
	0x2B	IRQ11 — eth0 (network)	
	0x2C	IRQ12 — PS/2 mouse	
	0x2D	IRQ13	
0x30–0x7F	0x2E	IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
	⋮	APIC vectors available to device drivers	
	0x80	system call vector (INT 0x80)	
0x81–0xEE	⋮	more APIC vectors available to device drivers	
	0xEF	local APIC timer	
0xF0–0xFF	⋮	symmetric multiprocessor (SMP) communication vectors	

```
static spinlock_t the_lock = SPIN_LOCK_UNLOCKED;
unsigned long flags;

/* start of critical section */
/* line 0 */ spin_lock_irqsave (&the_lock, flags);
/* line 1 */ old_head = head;
/* line 2 */ head = new_elt;
/* line 3 */ new_elt->next = old_head;
/* line 4 */ spin_unlock_irqrestore (&the_lock, flags);
/* end of critical section */

static spinlock_t the_lock = SPIN_LOCK_UNLOCKED;
unsigned long flags;

/* start of critical section */
/* line 0 */ asm volatile (" # local_irq_save macro implementation
    pushfl # save EFLAGS to stack
    popl %0 # pop EFLAGS into output 0
    cli # mask interrupts
": "g" (flags) /* output 0 is a general-purpose register */
: /* which should then be stored in flags */
: /* no inputs */
: "memory" /* see text */
);
spin_lock (&the_lock);
/* line 1 */ old_head = head;
/* line 2 */ head = new_elt;
/* line 3 */ new_elt->next = old_head;
/* line 4 */ spin_unlock (&the_lock);
asm volatile (" # local_irq_restore macro implementation
    pushl %0 # save input 0 to stack
    popl %0 # pop input 0 into EFLAGS
": /* no outputs */
: "g" (flags) /* input 0 is a general-purpose register */
: /* which should hold the value in flags */
: "memory", "cc" /* see text */
);
/* end of critical section */

/* Allocate statically and initialize to val. */
static __DECLARE_SEMAPHORE_GENERIC (name, val);
/* Allocate on stack and initialize to one. */
DECLARE_MUTEX (name);
/* Allocate on stack and initialize to zero. */
DECLARE_MUTEX_LOCKED (name);
```

initialization	
void spin_lock_init (spinlock_t* lock);	Initialize a dynamically-allocated spin lock.
basic lock and unlock functions	
void spin_lock (spinlock_t* lock);	Obtain a spin lock; call returns only when lock is obtained.
void spin_unlock (spinlock_t* lock);	Release a spin lock; must only be called on locks owned by caller.
miscellaneous testing functions	
int spin_is_locked (spinlock_t* lock);	Check if a spin lock is held. Returns 1 if held, 0 if available. Note that the lock may be claimed again before the caller can do anything!
int spin_trylock (spinlock_t* lock);	Make one attempt to obtain a lock. Returns 1 on success, 0 on failure.
void spin_unlock_wait (spinlock_t* lock);	Wait until a spin lock is available. Note that the lock may be claimed again before the caller can do anything!
lock and unlock with interrupt masking	
void spin_lock_irqsave (spinlock_t* lock, unsigned long* flags);	Save processor status in flags, mask interrupts, and obtain a spin lock; call returns only when lock is obtained.
void spin_unlock_irqrestore (spinlock_t* lock, unsigned long flags);	Release a spin lock, then set processor status to flags; must only be called on locks owned by caller.
void spin_lock_irq (spinlock_t* lock);	Mask interrupts and obtain a spin lock; call returns only when lock is obtained. Note that this version <i>does not preserve</i> the current value of the interrupt masking flag.
void spin_unlock_irq (spinlock_t* lock);	Release a spin lock, then enable interrupts; must only be called on locks owned by caller. Note that this version <i>does not restore</i> the previous value of the interrupt masking flag.

initialization	
void sema_init (struct semaphore* sem, int val);	Initialize a dynamically allocated semaphore to a value.
void init_MUTEX (struct semaphore* sem);	Initialize a dynamically allocated semaphore to the value one.
void init_MUTEX_LOCKED (struct semaphore* sem);	Initialize a dynamically allocated semaphore to the value zero.

down and up	
void down (struct semaphore* sem);	Wait on a semaphore; call returns only after success.
void up (struct semaphore* sem);	Signal a semaphore; must only be called by programs that have previously waited on the semaphore.

miscellaneous functions	
int down_interruptible (struct semaphore* sem);	Wait on a semaphore, but allow other programs to execute while waiting; call returns 0 on success or -EINTR on interruption.
int down_trylock (struct semaphore* sem);	Make one attempt to wait on a semaphore. Returns 0 on success, 1 on failure (<i>the opposite of the spin_trylock function!</i>).

```
void init_8259A(int auto_eoi)
{
    unsigned long flags;

    i8259A_auto_eoi = auto_eoi;

    spin_lock_irqsave(&i8259A_lock, flags);

    outb(0xff, 0x21); /* mask all of 8259A-1 */
    outb(0xff, 0xA1); /* mask all of 8259A-2 */

    /*
     * outb_p - this has to work on a wide range of PC hardware.
     */
    outb_p(0x11, 0x20); /* ICW1: select 8259A-1 init */
    outb_p(0x20 + 0, 0x21); /* ICW2: 8259A-1 IR0-7 mapped to 0x20-0x27 */
    outb_p(0x04, 0x21); /* 8259A-1 (the master) has a slave on IR2 */
    if (auto_eoi)
        outb_p(0x03, 0x21); /* master does Auto EOI */
    else
        outb_p(0x01, 0x21); /* master expects normal EOI */

    outb_p(0x11, 0xA0); /* ICW1: select 8259A-2 init */
    outb_p(0x20 + 8, 0xA1); /* ICW2: 8259A-2 IR0-7 mapped to 0x28-0x2f */
    outb_p(0x02, 0xA1); /* 8259A-2 is a slave on master's IR2 */
    outb_p(0x01, 0xA1); /* (slave's support for AEOI in flat mode
                        is to be investigated) */

    if (auto_eoi)
        /*
         * in AEOI mode we just have to mask the interrupt
         * when acking.
         */
        i8259A_irq_type.ack = disable_8259A_irq;
    else
        i8259A_irq_type.ack = mask_and_ack_8259A;

    udelay(100); /* wait for 8259A to initialize */

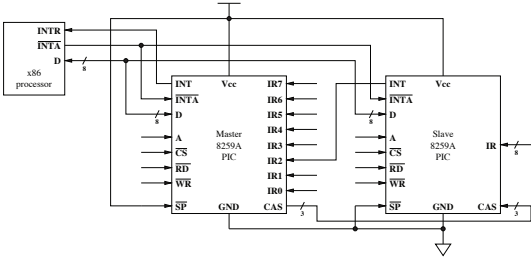
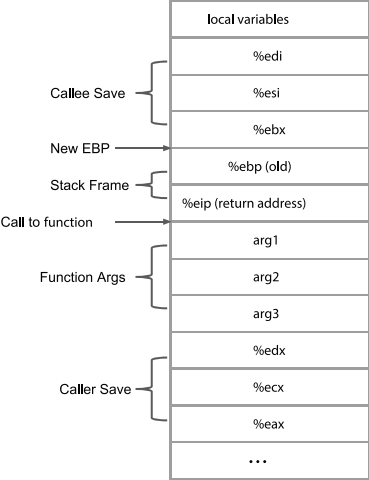
    outb(cached_21, 0x21); /* restore master IRQ mask */
    outb(cached_A1, 0xA1); /* restore slave IRQ mask */

    spin_unlock_irqrestore(&i8259A_lock, flags);
}
```

```
spin_lock:
movl 4(%esp), %eax
loop:
movl $1, %ecx
xchgl %ecx, (%eax)
cmpl $1, %ecx
je loop
ret
```

```
spin_unlock:
movl 4(%esp), %eax
movl $0, (%eax)
ret
```

function(arg1, arg2, arg3):



ADDR bus

